

## Q1. Model

### 1. Model

Describe the model architecture and how it works on text summarization.

The model I use in this homework is google/mt5-small, and its configuration is shown below.

```
{
  "_name_or_path": "google/mt5-small",
  "architectures": [
    "MT5ForConditionalGeneration"
  ],
  "d_ff": 1024,
  "d_kv": 64,
  "d_model": 512,
  "decoder_start_token_id": 0,
  "dense_act_fn": "gelu_new",
  "dropout_rate": 0.1,
  "eos_token_id": 1,
  "feed_forward_proj": "gated-gelu",
  "initializer_factor": 1.0,
  "is_encoder_decoder": true,
  "is_gated_act": true,
  "layer_norm_epsilon": 1e-06,
  "model_type": "mt5",
  "num_decoder_layers": 8,
  "num_heads": 6,
  "num_layers": 8,
  "pad_token_id": 0,
  "relative_attention_max_distance": 128,
  "relative_attention_num_buckets": 32,
  "tie_word_embeddings": false,
  "tokenizer_class": "T5Tokenizer",
  "torch_dtype": "float32",
  "transformers_version": "4.22.2",
  "use_cache": true,
  "vocab_size": 250100
}
```

MT5Model overrides T5Model, and T5Model is an encoder-decoder model. Since we're using summarization, the prefix "summarize:" will first be added. The model will first take the tokenized text input as the encoder's input. Then the output of the last layer

of the encoder will be the encoder hidden states of the decoder. The decoder will perform attention within inputs and encoder hidden states. Finally, the output of the decoder passes a linear layer, and the probabilities for the words are generated.

## 2. Preprocessing

Describe your preprocessing (e.g. tokenization, data cleaning and etc.)

- Tokenization:

MT5 applies T5 tokenizer. T5 tokenizer is based on SentencePiece, which implements subword units(e.g. byte pair encoding and unigram language model).

- Data cleaning:

The raw data needs to be preprocessed to make it accessible for the mt5 model. Text\_column and summary\_column will be set to "maintext" and "title" respectively. If the title is not given, then the target will be set to be [""] \* len.

## Q2. Training

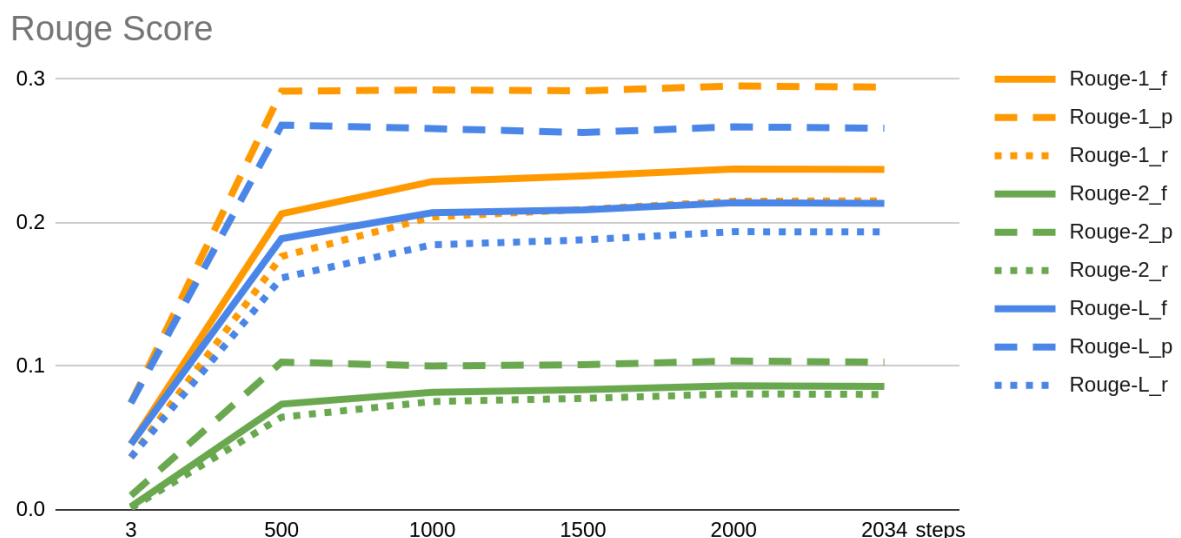
### 1. Hyperparameter

Describe your hyperparameter you use and how you decide it.

- fp16  
The memory requirement will be lessened during training.
- optim adafactor  
The optimizer will be adafactor instead of Adam to reduce the memory requirement of storing momentum.
- source\_prefix "summarize: "  
It is set to identify that it is doing summarization task.
- learning\_rate 1e-4  
A moderate learning rate that is not too low or too high.
- gradient\_accumulation\_steps 4
- eval\_accumulation\_steps 4
- per\_device\_eval\_batch\_size 4
- per\_device\_train\_batch\_size 4  
Gradient accumulation is applied to reduce the memory requirements for larger batch size.

### 2. Learning Curves

Plot the learning curves (ROUGE versus training steps)



## Q3. Generation Strategies

### 1. Strategies

Describe the detail of the following generation strategies:

- Greedy

In this strategy, each choice is greedily determined by the current most probable word. It does not consider previous word probability distribution into consideration.

- Beam Search

In this strategy,  $k$  is a hyper-parameter and  $k$  most probable sequences are being tracked. It will choose the best one among the  $k$  sequences at the end.

- Top-k Sampling

Instead of choosing the word with the most probable word directly, this strategy performs sampling from top  $k$  most probable words.

- Top-p Sampling

In this strategy, it adds most probable words into the sample set sequentially until the sum of their probabilities exceeds  $p$ . Then it samples the answer from the sample set.

- Temperature

Soft-max:

$$P(w_t) = \frac{e^{s_w}}{\sum_{w' \in V} e^{s_{w'}}}$$

Soft-max with temperature  $\tau$ :

$$P(w_t) = \frac{e^{s_w/\tau}}{\sum_{w' \in V} e^{s_{w'}/\tau}}$$

$\tau$  is high (high temperature): The distribution will be more unified.

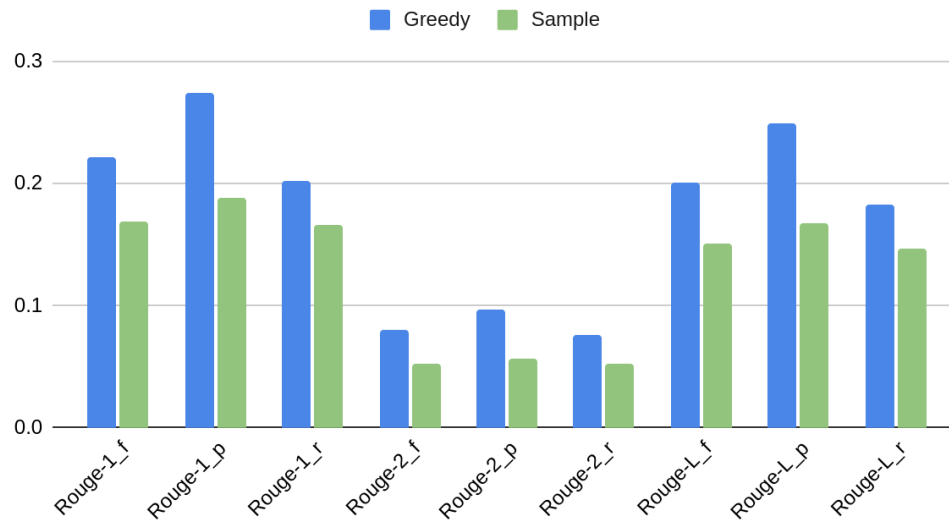
$\tau$  is low (low temperature): The distribution will be spikier.

### 2. Hyperparameters

Try at least 2 settings of each strategies and compare the result.

(a) Greedy

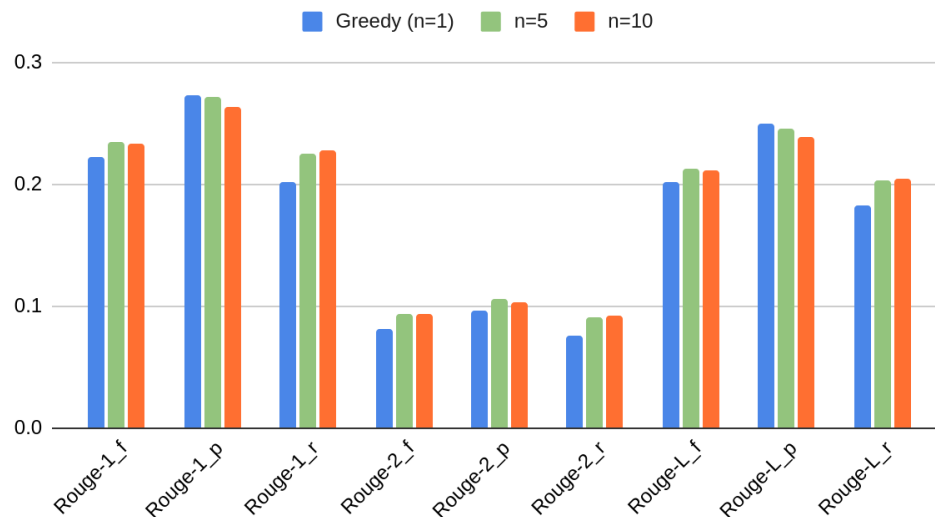
### Greedy and Sample



Obviously, greedy performs significantly better than sampling.

### (b) Beam Search

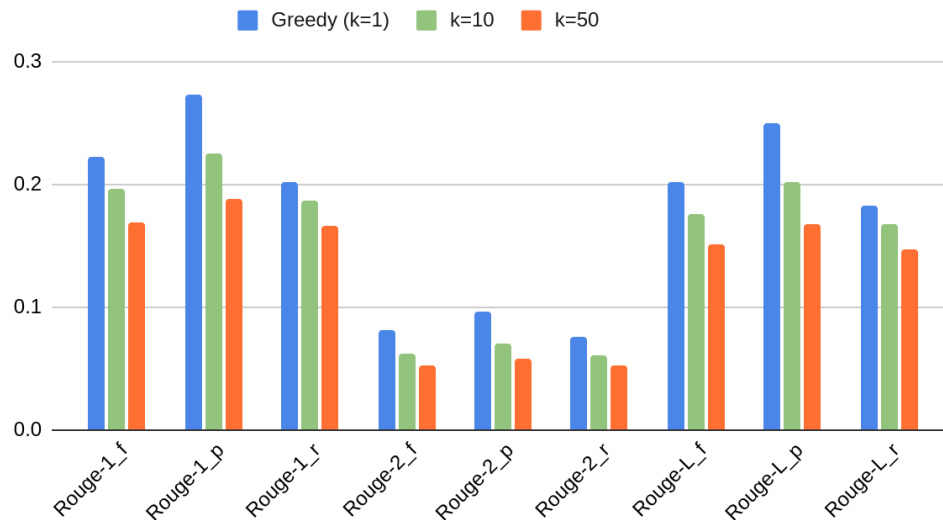
#### Beam Search



n=5 performs better than greedy in most cases. n=5 and n=10 have similar performance, but n=5 runs faster than n=10 does.

### (c) Top-k Sampling

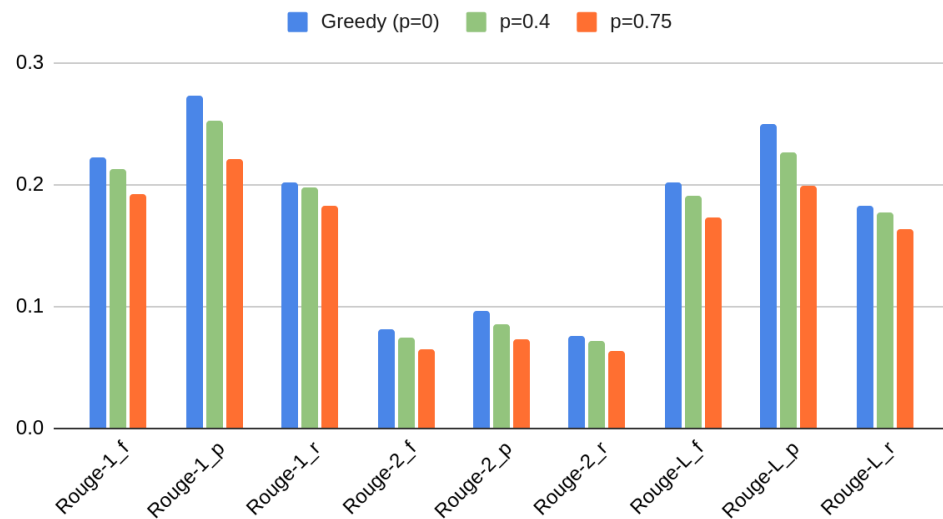
Top-k Sampling



As k becomes higher, it will worsen the performance from the experiment. Probably the selected k is too large.

(d) Top-p Sampling

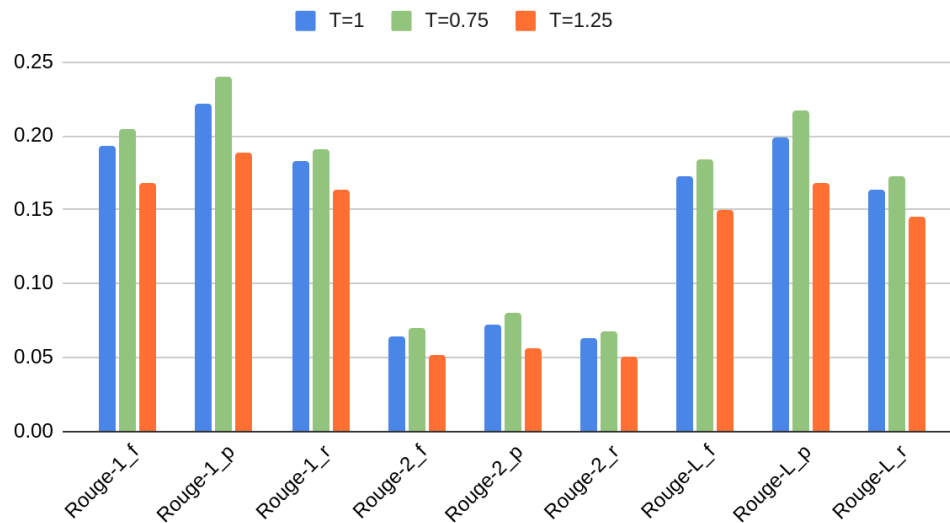
Top-p Sampling



Similar to top\_k, As p becomes higher, it will worsen the performance from the experiment. Probably the selected p is too large.

(e) Temperature

Temperature under  $p=0.75$



Under the settings of  $p=0.75$  and different temperatures,  $T=0.75$  has the best result out of the three settings.

What is your final generation strategy? (you can combine any of them)

My final generation strategy is numbeams=5.